

Sparsity-Adaptive Gated Eviction for Memory-Efficient Long-Context Inference of Large Language Models

Title Page — Applied Intelligence submission

Article type: Original Research
Author: Chang Tan (谭畅)
Position: Undergraduate Student
Affiliation: School of Economics and Management, Dalian University of Technology
Address: No. 2 Linggong Road, Ganjingzi District, Dalian 116024, Liaoning, China
E-mail: aquars43@foxmail.com
Corresponding author: The sole author is the corresponding author.
ORCID: (to be provided by the author, optional)

Note on the author name. The author's English name, Chang Tan (given name first, family name second), is confirmed by the author and is used consistently in the manuscript byline, this title page, and the cover letter. The Chinese characters, affiliation, address, and e-mail above are confirmed.

Declarations

Funding. The author did not receive support from any organization for the submitted work.

Competing interests. The author has no relevant financial or non-financial interests to disclose.

Data availability. The benchmark corpora used in this study (WikiText-2, PG19, LongBench, and the needle-in-a-haystack protocol) are publicly available through their original releases. The complete experimental record that reproduces every table and figure — per-seed values, per-task scores, and raw timing and memory counters — will be released in a public repository upon acceptance.

Code availability. The SAGE implementation and all configurations will be released under a permissive license upon acceptance.

Author contributions. The sole author performed all parts of the study and wrote the manuscript.

Acknowledgements. Not applicable; there is no acknowledgement content to declare.